1. 文章的 Pipeline（合成数据生成与蒸馏流程）
该研究构建了一个自动化的流水线，旨在将大模型（LLM）的能力蒸馏到 SLM 中，具体分为四个核心阶段

a. 正向查询生成（Positive Query Generation）： 从种子文档开始，利用大模型（如 GPT-4o）结合特定的“查询模式模板”（基于企业搜索中常见的作者、文件名、文件夹名等元数据组合）生成高质量、符合现实用户行为的企业搜索查询

b. 负样本挖掘（Negative Document Mining）： 针对生成的每一个合成查询，利用 BM25 算法从语料库中检索出排名前 k 的文档（排除原始文档），获取具有词汇或语义重叠的“硬负样本”

c. LLM 自动标注（LLM Labeling）： 使用教师模型（GPT-4o）根据特定的标注 Prompt，为上述查询-文档对分配 0（不相关）到 4（高度相关）的等级分数，形成合成训练集

d. SLM 标注员微调（SLM Labeler Fine-tuning）： 在 Phi-3.5 Mini Instruct 模型上进行多任务微调。除了合成的企业数据外，还引入了 INTERS、TREC-CAsT 和 MS MARCO 等公开数据集，以增强模型的泛化能力和对语义相关性的理解

2. contribution
a. 极高的性价比与吞吐量： 与 GPT-4o 相比，SLM 标注员的吞吐量提升了约 17 倍（达 873 RPM），而成本降低了约 19 倍，使得大规模离线评估变得经济可行

b. 性能媲美顶尖 LLM： 实验证明，经过微调的 SLM（Phi-3.5）在相关性标注的准确度（Pairwise Accuracy）和排序一致性（NDCG）上，不仅与教师模型 GPT-4o 持平，甚至在某些指标上略有超越

c. 高效的企业级标注方案： 提出了第一个专门针对企业搜索场景（结合元数据和混合检索信号）的 SLM 标注员微调方案

3. 这篇文章的思路

“大教小”的蒸馏理念： 大模型（LLM）擅长逻辑推理和少样本标注，但因其计算开销大，不适合大规模评估；小模型（SLM）在经过特定领域知识的蒸馏后，可以在特定任务（如相关性评分）上达到甚至超越通用大模型的表现

模式化的行为模拟： 通过分析用户在企业环境中的搜索行为模板（如“作者名+关键词”），可以引导 LLM 生成比随机文本更具行业真实感（Realism）的伪查询

多任务学习的稳健性： 虽然侧重企业数据，但同时引入大规模公开 Web 检索数据集进行辅助训练，能有效提升模型的泛化能力，防止在小规模合成数据上过拟合

arXiv:2601.03211v1 [cs.IR] 6 Jan 2026

FINE-TUNING SMALL LANGUAGE MODELS AS EFFICIENT ENTERPRISE SEARCH RELEVANCE LABELERS

Yue Kang[1*†] **Zhuoyi Huang**[1*] **Benji Schussheim**[1] **Diana Licon**[1] **Dina Atia**[1]
Shixing Cao[1] **Jacob Danovitch**[1†] **Kunho Kim**[1] **Billy Norcilien**[1] **Jonah Karpman**[1]
Mahmound Sayed[1] **Mike Taylor**[1] **Tao Sun**[1] **Pavel Metrikov**[1] **Vipul Agarwal**[2‡]
Chris Quirk[1] **Ye-Yi Wang**[1] **Nick Craswell**[1] **Irene Shaffer**[1] **Tianwei Chen**[1†]
Sulaiman Vesal[1†] **Soundar Srinivasan**[1†]
[1]Microsoft [2]Amazon

ABSTRACT

In enterprise search, building high-quality datasets at scale remains a central challenge due to the difficulty of acquiring labeled data. To resolve this challenge, we propose an efficient approach to fine-tune small language models (SLMs) for accurate relevance labeling, enabling high-throughput, domain-specific labeling comparable or even better in quality to that of state-of-the-art large language models (LLMs). To overcome the lack of high-quality and accessible datasets in the enterprise domain, our method leverages on synthetic data generation. Specifically, we employ an LLM to synthesize realistic enterprise queries from a seed document, apply BM25 to retrieve hard negatives, and use an a teacher LLM to assign relevance scores. The resulting dataset is then distilled into an SLM, producing a compact relevance labeler. We evaluate our approach on a high-quality benchmark consisting of 923 enterprise query–document pairs annotated by trained human annotators, and show that the distilled SLM achieves agreement with human judgments on par with or better than the teacher LLM. Furthermore, our fine-tuned labeler substantially improves throughput, achieving a $17\times$ increase while also being $19\times$ more cost-effective. This approach enables scalable and cost-effective relevance labeling for enterprise-scale retrieval applications, supporting rapid offline evaluation and iteration in real-world settings.

1 Introduction

As AI-driven search technology advances, Large Language Models (LLMs) have demonstrated remarkable capabilities in search and relevance labeling [1, 2]. However, enterprise search presents unique challenges that go beyond those of traditional web search [3]. Enterprise search refers to retrieving information from an organization's heterogeneous internal repositories, such as emails, chats, documents, and knowledge bases, where users issue both keyword-based queries (e.g., searching by file name or author) and semantic queries (e.g., asking natural language questions about documents). Unlike web search queries, enterprise queries are often ambiguous, persona-specific, and heavily context-dependent, requiring systems to interpret user intent while accounting for content, metadata, and both personal and organizational context [4]. For example, while the query "Juno release date" has a clear intent in web search relating to the 2007 film [4], the same query in enterprise could have multiple intents, e.g. looking for files containing the release date of a project named Juno or for emails from a colleague named Juno. The differences between enterprise search queries, the focus of our work, versus web search queries are also presented in Table 1.

Consequently, labeling relevance in enterprise search is inherently more complex than in web search. Moreover, there is a lack of publicly available, high-quality enterprise search datasets that align with real-world data structures and distributions, limiting the development of enterprise search methods. In contrast, web search has benefited from large-

[*]Both authors have equal contribution.
[†]Corresponding authors: {yuekang,jdanovitch,tianweichen,svesal,sosrini}@microsoft.com
[‡]This work was conducted while employed at Microsoft.
[4]https://en.wikipedia.org/wiki/Juno_(film)

scale public benchmarks such as MS MARCO [5] and TREC-CAsT [6], which have driven rapid progress in retrieval and relevance labeling. In addition, personalization and contextualization are essential, as relevance depends on both textual metadata and user networks. Finally, enterprise queries are often underspecified, domain-specific, and long-tail, making ranking evaluation particularly challenging due to sparse user engagement and limited human-labeled data. These challenges underscore the value of LLMs for relevance labeling [7]. However, while LLM-based judgments [8] have become a standard approach for enterprise relevance labeling, they come with significant drawbacks: LLMs are computationally expensive and have limited throughput, making them inefficient for large-scale labeling tasks. To address these limitations, Small Language Models (SLMs) [9] offer a promising alternative. With domain-specific fine-tuning on curated synthetic data, SLMs can achieve decent performance comparable to LLMs while significantly improving throughput and reducing costs.

Recent studies have investigated training SLMs for labeling web search queries [10], but enterprise scenarios demand broader coverage of hybrid queries that combine lexical and semantic signals, a setting that remains largely unexplored. To address this gap, we propose a novel synthetic data generation pipeline that fine-tunes SLMs to label enterprise search queries at scale. Our approach leverages enterprise documents to seed the generation of realistic user queries, applies BM25 [11] to retrieve challenging negatives, and employs LLM-based judgments to create graded relevance labels. These synthetic datasets are then distilled into an SLM, yielding an efficient and cost-effective relevance labeler capable of handling diverse enterprise query types. This method enhances the offline evaluation process for ranking models, offering a scalable alternative to conventional LLM-based labelers and human judges. Our main contributions are as follows:

- We propose a novel synthetic data generation pipeline that leverages BM25 [11] and LLM-based labeling to create high-quality, domain-specific, and realistic training data for enterprise search relevance labeling, addressing existing dataset limitations.
- We showcase the effective fine-tuning of an SLM, Phi-3.5 Mini Instruct [12], achieving performance comparable to its teacher LLM (GPT-4o) while significantly improving throughput and reducing costs.

These facts demonstrate that the fine-tuned SLM enables cost-effective, production-scale offline ranking evaluation, offering a practical alternative to LLM-based or human labeling.

2 Related Work

Large and Small Language Models The rapid progress of generative language models has transformed natural language processing (NLP) and reshaped modern information retrieval (IR), enabling systems to perform advanced reasoning and comprehension across a wide spectrum of applications [13, 14, 15]. Early innovations such as ELMo [16] and BERT [17] demonstrated that large-scale pretraining on raw text could capture rich contextual dependencies, and later refinements like ELECTRA [18] and DeBERTa [19] introduced more efficient and accurate pretraining strategies. The introduction of sequence-to-sequence architectures, such as T5 [20], has enabled a unified treatment of diverse IR tasks, ranging from query rewriting to passage ranking. Embedding-based retrieval methods, such as ColBERT [21], further highlighted how encoders can surpass lexical matchers by capturing nuanced semantic signals. In parallel, decoder-only models such as GPT [22, 23] marked the advent of large-scale generative systems and rapidly extended applications to document ranking, conversational search, and summarization [24], ultimately establishing themselves as state-of-the-art engines for modern IR.

More recently, attention has shifted from larger models to compact yet competitive small language models (SLMs). Families such as Gemma [25], Llama [26], and Phi series [9, 12] demonstrate that models with a few billion parameters can offer strong generalization and reasoning capacity while drastically lowering inference cost. While SLMs typically fall short of frontier LLMs in absolute performance, their speed, lower memory footprint, and reduced inference cost make them highly appealing for retrieval and labeling tasks where throughput and scalability are critical [27]. This trade-off, accepting slightly lower peak accuracy in exchange for efficiency, allows SLMs to serve as practical complements to LLMs, enabling high-volume offline evaluation and rapid experimentation without the burden of massive GPU resources.

Relevance Labeling with Language Models Relevance assessment has long been a cornerstone of IR, traditionally dependent on human annotations along with heuristic models such as BM25 [11]. While accurate, these methods were costly and limited in scale. More recently, LLMs have been explored as alternatives to replicate or augment human judgments. [28] introduced one-shot labeling for automatic relevance estimation, showing that LLMs could generate labels efficiently. [29] examined the use of LLMs to fill relevance judgment holes, and later [8] showed that LLMs could accurately predict searcher preferences, aligning closely with user behavior, while [30] proposed criteria-based LLM judgments to improve label consistency across dimensions such as topicality and coverage.

Query Type	Enterprise	Web
Description	Queries leverage context and internal entities along with semantic understanding	Queries rely on common knowledge and general semantics
Example	iris project date	structure of the iris
Key Difference	Further require enterprise-specific knowledge (e.g., "iris" = colleague/project)	Use global/common knowledge (e.g., "iris" = part of the eye)

Table 1: Enterprise vs. web search: both can be semantic, but enterprise queries depend on domain-specific knowledge.

At the same time, fine-tuning techniques such as LoRA [31] have made it feasible to adapt LLMs to domain-specific ranking tasks [32]. More efficient distillation-based methods have also emerged: [10] demonstrated that fine-tuned small LLMs could improve the quality of augmented relevance datasets, highlighting the promise of post-training SLMs for scalable evaluation. RRADistill [7] distilled large models into SLMs for re-ranking long-tail queries, while Rank1 [33] leveraged reasoning-based LLMs (e.g., OpenAI's o1 [34], DeepSeek's R1 [35]) and MS MARCO-derived traces to train smaller, explainable, high-performing re-rankers. Collectively, these approaches highlight the promise of synthetic data pipelines and distillation for building compact yet effective retrieval models.

Despite this progress, the exploration of SLMs for relevance labeling has been largely limited to open-domain, semantically driven retrieval tasks [10]. While these works showed that compact models can enhance dataset quality or approximate semantic relevance, enterprise search poses distinct challenges. Queries in enterprise environments often combine sparse keyword matches with semantic intent, and document collections span heterogeneous domains such as Teams chats, emails, and files. Current industrial practice continues to rely on frontier-scale LLMs, combined with careful prompt engineering and chain-of-thought reasoning [36], to handle this complexity [8]. To the best of our knowledge, no prior work has examined whether fine-tuned SLMs can deliver competitive and resource-efficient labels in this setting. Our work addresses this gap by adapting Phi-3.5-mini for enterprise relevance labeling, demonstrating that carefully curated training datasets and distillation pipelines can make SLMs a faster, more efficient, and practical alternative for large-scale experimentation.

Synthetic data generation with LLMs to create queries has proven effective for model training in low-resource settings [37]. Retrieval approaches fine-tuned solely on synthetic data have outperformed strong baselines like BM25 and other self-supervised dense retrieval methods across various domain-specific IR tasks [38]. Approaches such as Promptagator [39] and InPars v2 [40] have enabled the development of open-source rankers that achieve state-of-the-art results on public benchmarks, while DUQGen [41] explored unsupervised domain adaptation through clustering and probabilistic sampling to diversify synthetic queries.

Building on these innovations, we focus on the practical deployment of a Small Language Model (SLM) labeler for enterprise search and retrieval. Specifically, we fine-tune Phi-3.5 Mini [9] using domain-specific data generated through an LLM-driven query and label pipeline, customized via instruction tuning on synthetic enterprise-style queries and documents. We selected Phi-3.5 Mini and GPT-4o as benchmark models, given their proven performance in recent retrieval and reasoning studies [24], and because they represent strong exemplars of their respective model classes (compact SLMs vs. frontier-scale LLMs). While we did not include the results with newer models such as GPT-5 or alternative SLM families, we believe our findings to generalize, since the core advantage comes from the data generation and distillation pipeline rather than any model-specific property. Our resulting SLM labeler achieves performance comparable to frontier LLMs while delivering substantial improvements in efficiency and scalability, making it well-suited for large-scale enterprise relevance evaluation. Notably, traditional approaches to search relevance labeling, such as manual annotation or direct reliance on frontier LLMs, present significant challenges. Manual annotation is often infeasible in real-world enterprise search due to strict privacy policies and the substantial effort required to go through large-scale datasets. LLMs-based labeling, while capable of producing high-quality labels, is computationally expensive and slow, making it unsuitable for large-scale deployment.

2.1 Problem Definition

Given a user query q, the objective is to determine the relevance of a set of documents $D = \{d_1, d_2, ..., d_n\}$ to the query q. This can be formalized as learning a relevance function $r(q, d_i)$ that assigns a relevance score to each query-document pair, (q, d_i), where $d_i \in D$. We define the problem as follows:

Input:

- A query $q \in \mathcal{Q}$, where $\mathcal{Q}$ is the space consisting of all user enterprise queries.
- A set of documents $D = \{d_1, d_2, ..., d_n\}$, where $d_i \in \mathcal{D}$ and $\mathcal{D}$ are the space of documents.

Dataset	# Samples	Description	Usage
Eyes-on enterprise documents	1,500	A proprietary collection of internal documents with associated metadata, curated under an eyes-on review setting.	**Supervised Fine-tuning:** Used to generate synthetic query–doc–label triplets and fine-tune the model on them.
INTERS [42]	∼250,000 samples on 20 different tasks	Multi-task tuning dataset covering multiple tasks, including query, document, and query-document relationship understanding.	**Supervised Fine-tuning:** Used for multi-task tuning. (Green box in Figure 3)
TREC-CAsT [6]	4,000	Public dataset for Conversational Information Seeking (CIS) with passage-query-human label triplets.	**Supervised Fine-tuning:** Used ∼4k samples for instruct tuning. (Blue box in Figure 3)
MS MARCO Passage [5]	Over 400,000 passages	Open dataset for information retrieval, containing passage-query pairs and some of them have human annotations of scale 0-3.	**Supervised Fine-tuning:** Used GPT-4o to generate synthetic queries for instruct tuning. (Purple box in Figure 3)
Human-labeled enterprise search dataset	923	A proprietary dataset of enterprise query–document pairs, annotated by trained human labelers for relevance.	**Evaluation:** Used as the gold-standard evaluation dataset for final comparison.

Table 2: Datasets used for fine-tuning and evaluation

Output:

- A relevance score $r(q, d_i) \in \mathcal{R}$ for each (q, d_i) pair, where $\mathcal{R}$ is the space of relevance scores. In this paper, we adopt an ordinal 0–4 scale, with 0 indicating irrelevance and 4 indicating high relevance. Scores $r(q, d_i)$ are designed to approximate human judgments and align with the quality of state-of-the-art LLMs such as GPT-4o.

Our goal is to develop a practical and scalable approach for enterprise search relevance labeling, enabling rapid and cost-effective evaluation of various IR models while ensuring high-quality and consistent supervision in real-world information filtering systems.

2.2 Datasets

Table 2 provides detailed descriptions of the datasets and specifies their roles across the different phases from model training to evaluation. Specifically, we use four datasets for training and one benchmark dataset for final evaluation. We obtain 1,500 proprietary enterprise documents curated under an eyes-on review setting for synthetic data generation and fine-tuning, while we believe a similar dataset could be obtained from any document platform. Since no public enterprise search dataset is available for evaluation, we curate an internal benchmark consisting of 923 high-quality enterprise query–document pairs annotated by trained human annotators.

2.3 Evaluation Metrics

To assess the effectiveness of our SLM relevance labelers against baseline labelers, we employ two complementary evaluation metrics: Normalized Discounted Cumulative Gain (NDCG) [43] and Pairwise Accuracy. NDCG is a widely adopted metric in information retrieval that evaluates the quality of ranked lists with respect to graded relevance, while Pairwise Accuracy measures how well the model preserves the correct relative ordering of documents under the same query. A higher level of both metrics indicates that the model's predictions are more consistent with the intended relevance ordering. After computing both metrics at the query level, we report their average values across different queries.

1. Normalized Discounted Cumulative Gain (NDCG): We adopt full Normalized Discounted Cumulative Gain (NDCG) to capture model performance across entire rankings under each query. NDCG measures ranking quality by assigning higher gains to more relevant documents and discounting them logarithmically by their rank positions, ensuring that placing relevant items earlier in the list is rewarded more. Unlike truncated variants (e.g., NDCG@k), full NDCG accounts for all retrieved documents for each query, offering a comprehensive and widely recognized evaluation of ranking effectiveness.

2. Pairwise Accuracy Metric (Accuracy): This metric measures how consistently the model preserves the correct relative ordering of documents compared to the ground-truth labels. For each query, we exhaustively consider all possible pairs of documents (A, B) and check whether the model's prediction of their relative relevance agrees with the ground truth. For example, if both the model and the ground truth indicate that A is more relevant than B ($A > B$), then this pair is counted as correctly aligned. We compute the proportion of correctly aligned pairs across all possible combinations for a query, resulting in a score between 0 and 1, and then average these scores across queries. Table 3 illustrates the full computation matrix. A cell value of 1 indicates perfect agreement between the model and the ground truth for that comparison, while 0 indicates a mismatch.

Model Output	$A<B$	$A = B$	$A>B$
Model says $A<B$	1	0	0
Model says $A = B$	0	1	0
Model says $A>B$	0	0	1

Table 3: Accuracy metric details. Here: 1: The model's prediction exactly matches the baseline label, 0: The model's prediction does not match the baseline label.

3. Request Per Minute (RPM): We also report RPM, which is a metric used to measure the rate at which our SLM labeler processes requests on declared hardware. It is commonly used to monitor system performance, rate limits, and scalability.

$$\text{RPM} = \frac{\text{Number of total requests for labeling}}{\text{Time elapsed for the labeling in minutes}}$$

3 Methods

Unlike web search, where large-scale open-source datasets such as MS MARCO [5] and TREC-CAsT [6] provide high-quality human-labeled query–document pairs, there is no publicly available dataset for enterprise search that captures its unique characteristics. While some enterprise systems may collect query logs, this is not always the case, and oftentimes these contain sensitive, confidential user information which makes it untenable to collect human relevance judgements. Moreover, interaction data such as click logs are inherently biased, as they primarily reflect user engagement rather than true relevance [44]. For example, users often click on higher-ranked results regardless of their quality [45]. These limitations underscore the need for an alternative approach that enables effective training while addressing the absence of feasible training data.

To this end, we propose two approaches for constructing fine-tuning datasets for enterprise relevance labeling without relying on large-scale query logs. Instead, both methods require only a curated set of seed documents. As a result, our methodology is fully automatic and inherently privacy-preserving, making it well suited for enterprise settings where user queries are sensitive, scarce, or unavailable. Specifically, both approaches are based on the idea of generating synthetic queries from the aforementioned seed documents. To generate synthetic queries for fine-tuning our SLM labeler, we investigate two different approaches. The first approach fine-tunes an auxiliary SLM to generate queries given a seed document and a predefined relevance score, enabling synthetic query creation through random sampling of document–score pairs. These generated queries are then used for training the SLM labeler. The second approach, in contrast, involves creating synthetic query-document-score triplets using LLM-based query generation, BM25-based negative document mining, and LLM-assisted labeling, followed by fine-tuning the SLM labeler on these triplets. In practice, we found that the first method struggled to produce high-quality and diverse queries, particularly for keyword-based enterprise search, limiting its utility. Consequently, we focus primarily on the second approach, which demonstrated more consistent and reliable performance. We describe both methods in detail below.

3.1 Fine-tune SLMs as Query Generators

We first fine-tuned the Phi-3.5 Mini Instruct model as a query generator, as this approach enables faster large-scale query generation with controlled relevance levels while reducing the reliance on expensive LLM inference. To the best of our knowledge, there is no prior work on fine-tuning an SLM for enterprise search query generation. Additionally, we previously had fine-tuned the same SLM for semantic web query generation, conditioning it on both the context and a predefined relevance score. This method showed promising results in key evaluation metrics, which motivated us to extend the method to the enterprise setting. The key differences between enterprise search queries, which are the focus of this work, and web search queries are presented in Section 1.

3.1.1 Experimental Setup

To generate synthetic queries for fine-tuning our SLM query generator, we utilized GPT-4o to produce both positive and negative enterprise search queries using a carefully designed prompt with explicit instructions for diverse example generation. We then fine-tuned Phi-3.5 Mini Instruct on this synthetic dataset under multiple training configurations aimed at enabling the SLM to generate enterprise-style queries conditioned on document metadata and target relevance levels. However, when we evaluated the fine-tuned model by comparing its outputs against GPT-4o gold-standard judgments (LLM-as-a-judge), we observed that it achieved only 60.8% binary relevance accuracy, i.e. only 60.8% of its generated queries matched the intended relevance class (relevant or irrelevant) evaluated by GPT-4o. A deeper manual analysis revealed a strong bias toward generating positive or plausible-sounding queries, even when the prompt

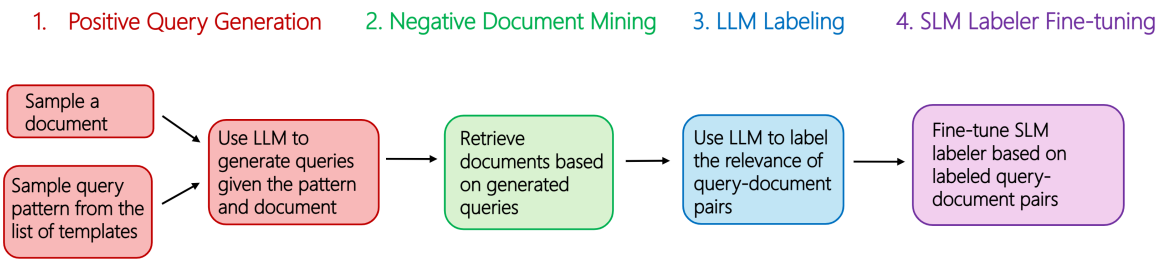

Figure 1: Overall model architecture for our proposed solution.

Example:
File Content: (omitted here due to the space limit)
Metadata: File name: DB Moves budget plan decent.one | Author: John Smith | Parent Folder: Open Projects opportunity | Extension: OneNote | Title: DB Moves budget plan decent
Input pattern: <unk>, 'File Name', 'Author', 'Parent Folder'

Result from 1st LLM call (generation):
Keywords: Cloud, Budget, Module, Security, Container
Queries: Cloud *DB Moves budget* John Open Projects, Module *DB Moves budget* John Open Projects, Security *DB Moves budget* John Open Projects (It uses same segments of different metadata all the time!)

Result from 2nd LLM call (revision for variability):
Revised Queries: Cloud *DB Moves budget* John Open Projects, Module *DB budget plan* John Smith Projects, Security *Moves descent* John Projects

Figure 2: An illustration of our synthetic query generation procedure.

explicitly instructed the model to produce irrelevant (negative) ones. As an illustrative instance, we fabricate a sample document to avoid privacy concerns and evaluate our SLM query generator:

- Content Summary: A document about how to add a page in Word.
- Keywords: page, Word, add
- Metadata: File name: ”AddPage.docx”, Author: Lisa Morrison, Title: ”AddPage”, File type: ”docx”, Parent folder: ”Word Tutorial”

By instructing the model to generate negative queries, we get: ”Lisa tutorial Docs”, ”Add Page” and ”Lisa Loop page”, but all these three queries are relevant with the document.

We attribute this limitation to two main factors:

- Positivity bias in pretraining: Instruction-tuned SLMs like Phi-3.5 Mini are optimized to be helpful and cooperative. When asked to generate a query ”for a document that should not be retrieved,” the model may still attempt to generate something that sounds relevant or plausible, rather than something truly irrelevant even after fine-tuning.
- Limited model capacity: As a small model, Phi-3.5 Mini may struggle to internalize nuanced instructions like generating “key-word-based queries that appear plausible but are not relevant,” especially in enterprise contexts where relevance boundaries can be subtle.

Overall, while the fine-tuned model demonstrated a decent ability to generate queries aligned with semantic information from documents, it consistently struggled with the keyword-based nature of enterprise search queries. Unlike traditional web searches, where queries often reflect user intent and semantic understanding, enterprise search further relies on exact keyword matching, file metadata, and structured document attributes. As a result, fine-tuning on SLMs fails to generate meaningful hard negative and diverse queries, and ultimately does not work well under real-world enterprise search applications.

3.2 Using LLMs and Retrievers for Query Generation

This alternative approach constitutes the core methodology of our work, integrating synthetic data generation with information retrieval techniques. Specifically, we first used GPT-4o to generate and refine synthetic enterprise search

queries from documents. Next, we applied ranking algorithms such as BM25 to mine plausible negative documents. With the resulting query–document pairs, we then prompted GPT-4o to assign graded relevance labels using a carefully designed labeling prompt. Finally, we fine-tuned Phi-3.5 Mini Instruct on these query–document–label triplets.

The overall workflow of this synthetic query generation pipeline is illustrated in Figure 1. It consists of four major components: (1) positive query generation, (2) negative document mining, (3) LLM-based labeling, and (4) SLM labeler fine-tuning. We describe each component in detail below.

Box 3.1: Positive Query Generation Prompt

Role
You are an assistant generating **keywords** and **user search queries** from file metadata and content.

Input
– A structured string of file metadata parts.
– Text content of the file.
– Metadata may contain content placeholders.

Task
Step 1: Keyword Extraction
– Extract up to 6 relevant single-word keywords from the content.
– Exclude stop words.
– If content is empty, return an empty list.

Step 2: Query Generation
– Use the metadata string in exact order to form natural queries.
– Replace each content placeholder with keyword.
– Generate 3 distinct queries without reordering parts.

Examples
– *Few-shot examples shown here...*

Output Format
```
Keywords:  k1, k2, k3, ...
Queries:  q1, q2, q3
```

1. Positive query generation: We began with a document dataset and a query pattern template table. In practice, such templates can be constructed by analyzing how users typically form queries in enterprise search. For example, a pattern like <author name><file name>might yield the query "Lisa budget report."

In our work, we derived an internal template table containing a ranked list of common query patterns (e.g., <folder name><document type><keyword>) along with their frequencies in actual search traffic. However, even without access to such data, synthetic template tables can be generated by systematically enumerating combinations of document metadata fields and keywords, which we found to produce realistic and diverse query structures. The patterns were constructed via entity resolution of query text, ensuring that each segment corresponded to an entity in the file metadata. In addition, templates may include placeholders for content-derived information, where document keywords extracted from the text can be inserted. As such, they accurately reflect how users typically search for documents in workplace settings.

To ensure our synthetic queries mirrored realistic user behavior, we employed weighted sampling over the query pattern table, selecting each pattern with probability proportional to its observed frequency (with replacement). Thus, more frequent patterns were sampled more often, preserving the distributional properties of real enterprise queries. Each sampled pattern served as a template, into which relevant document metadata (e.g., author, title, file type, parent folder) and extracted keywords were dynamically inserted.

Once a pattern-document pair was selected, we leveraged GPT-4o to generate three distinct enterprise queries for that pair. Each query generation step also included a lightweight keyword summarization phase, where the LLM extracted key thematic terms from the document to incorporate into the queries. This helped ground the queries generated in the actual content of the document. The structure of the used prompt is given in Box 3.1. Due to confidentiality constraints, we cannot disclose the full prompt templates. Instead, we provide a high-level outline to illustrate the general structure. The precise wording is not central to our contribution; rather, our key finding is that fine-tuned small language models can effectively perform enterprise relevance labeling across variations in prompt phrasing.

For more complex patterns, those with more components or rare metadata combinations, we observed that the initial queries sometimes lacked variation. For example, if the sampled query template was long, the LLM tended to repeatedly select the same portion of each component, even when explicitly instructed to vary them. To address this limitation, we added an additional LLM-based query revision stage that rewrote the queries for diversity, fluency, and coverage, thereby producing a richer set of realistic queries while preserving their core semantic relevance. This query revision prompt is outlined in Box 3.2. To illustrate how the two prompts work in sequence, Figure 2 shows a example (with fabricated metadata). The synthetic query generation prompt produces three reasonable queries, but they repeatedly use the same metadata segments, limiting diversity. The refinement step then rewrites these queries under strict rules, yielding more diverse outputs.

This process resulted in a high-quality positive synthetic query-document pairs pool and yielded an average of three diverse, metadata-aligned queries per document-template pair.

Box 3.2: Positive Query Revision Prompt

Role
You are an assistant validating and revising **user-generated search queries**.

Input
– Metadata string (with possible keyword placeholders)
– Keywords list
– Three generated queries

Task
Step 1: Validation
– Ensure queries follow metadata order and structure.
– Replace each content placeholder with keyword.
– Avoid redundancy; keep queries short and natural.

Step 2: Modification
– Revise queries so that each metadata part is phrased differently across the three queries.

Examples
– *Few-shot examples shown here...*

Output Format
`Revised Queries: query1, query2, query3`

2. Negative document mining: To construct a balanced and high-quality training set for relevance modeling, we employ the Okapi BM25 retrieval algorithm [46] to mine negative documents for each synthetic query. BM25 is a strong lexical-based baseline widely used in traditional information retrieval, and it allows us to simulate a realistic retrieval environment in the absence of human click logs or explicit negative labels.

For each synthetically generated query, we use BM25 to retrieve the top-k documents from the corpus, excluding the original source document from which the query was generated. Importantly, we do not assume that the source document is the only relevant document, nor do we require all retrieved documents to be strictly non-relevant. Instead, the retrieved candidates are treated as plausible retrieval results that may span a range of relevance levels. In other words, our goal at this stage is not to pre-classify retrieved documents as relevant or non-relevant, but to collect a realistic spectrum of relevance levels, with final relevance judgments deferred to the subsequent LLM-based annotation stage.

In practice, BM25 tends to return documents with varying degrees of lexical and semantic overlap with the query. Some retrieved documents may be highly relevant (e.g., discussing the same entity or topic), while others are only partially related or largely irrelevant. This behavior aligns well with our objective to construct training data that covers a diverse spectrum of relevance scores. As a result, the retrieved documents naturally form a mixture of hard negatives, weak positives, and near-miss cases, reflecting the ambiguity commonly observed in real-world enterprise search. Such diversity is crucial for teaching the model to distinguish truly relevant documents from those that are only topically or semantically adjacent. For example, given the query "Lisa budget report" (derived from a document authored by Lisa Morrison titled "2023 Budget Report" mentioned in Subsection 3.1), BM25 may retrieve another document such as "Lisa travel plan" or "2022 Budget Draft". These results are lexically similar and thus plausible retrievals, but they are not the intended match, making them suitable hard negatives for training.

In our experiments, we set $k = 4$, which empirically yields a uniform distribution of relevance labels across levels 0–4 after LLM-based annotation. This balance is desirable for training a robust relevance labeler, as it prevents

over-representation of trivially irrelevant examples while encouraging the model to learn subtle distinctions between closely related documents. In practice, k can be adjusted dynamically based on corpus size and diversity to maintain balanced coverage across relevance levels. For example, if the LLM-based annotation results in an over-representation of higher relevance scores, we can increase k to retrieve additional candidate documents and annotate only the newly introduced query–document pairs, while retaining previously labeled pairs, thereby incrementally rebalancing the dataset without reprocessing the entire corpus.

3. LLM labeling: The next stage in our data pipeline is LLM-based labeling, where GPT-4o assigns relevance scores to each query–document pair. This step is critical for distilling the retrieval capabilities of LLMs into our SLM labeler. We adopted a generic labeling prompt template, which instructs GPT-4o to output a discrete relevance score, ranging from 0 (irrelevant) to 4 (highly relevant), for each query–document pair, as outlined in Box 3.3.

To ensure accuracy and consistency, we introduced an additional quality-control round after the initial labeling. In this phase, we automatically checked GPT-4o outputs for formatting issues and semantic inconsistencies. Out of roughly 24,000 samples, only 15 cases contained invalid or incomplete outputs (e.g., missing scores), which were automatically re-labeled. Furthermore, we implemented a post-labeling filtering mechanism to preserve dataset quality. Specifically, if a synthetically generated positive query from Step 1 received an unexpectedly low score (0 or 1), we re-labeled it with GPT-4o using the same prompt. If the low score persisted, we retained it as a valid negative case; otherwise, we discarded the query. This filtering process prevents the model from learning contradictory signals and enforces a consistent training distribution.

After completing these four steps, we obtained a high-quality set of synthetic query–document–label triplets, which served as the core training dataset for fine-tuning our SLM labeler.

4. SLM labeler fine-tuning: Finally, we fine-tuned the Phi-3.5 Mini Instruct model using the labeled query-document pairs. To optimize the instruct fine-tuning process, we adapted the same prompt template for LLM labeling in the previous step after removing the instruction for the model to generate explanations alongside the predicted relevance score. Through empirical analysis, we found that including explanations, though useful for interpretability, led to inconsistent training signals and introduced unnecessary verbosity in the trained SLM's output. To further improve the generalization ability of the fine-tuned small language model, we employed multi-task tuning [47]. As illustrated in Figure 3, this phase incorporated several additional datasets to expose the model to a broader range of query-document relevance scenarios. Specifically, we include:

- INTERS [42]: It consists of data sets for 19 different small tasks about query and document understanding. We used it for multi-task tuning to enhance models' generalizability.

- TREC-CAsT [6]: This dataset consists of semantic web query-passage pairs with human labels on scale of 0-4.

- MS MARCO passages [5]: It consists of over 400K passages from various sources. We use GPT-4o to generate sematic queries under 0-4 relevance scale under randomly selected passages based on a chain-of-thought idea.

Although our primary goal is enterprise query labeling, we deliberately incorporated semantic web query datasets from the public domain. We believe this broader supervision helps the model develop a stronger understanding of query–document relevance beyond strictly keyword-matching setting, thereby improving both generalization and robustness. Specifically, as shown in Figure 3, we first trained the Phi-3.5 Mini Instruct on INTERS for multi-task tuning to enhance its robustness and overall capabilities, and then we continuously fine-tuned the trained model on a combination of three datasets we mentioned above: synthetic enterprise query-document-label triplets, TREC-CAsT with human labels, and MS MARCO passages with synthetic queries.

We train the model for 2 epochs using a maximum sequence length of 4096. The effective batch size is 32, achieved by setting the per-device batch size to 4 and using 8 gradient accumulation steps. Training logs are recorded every 40 steps. Evaluation is performed every 80 steps using a separate human-labelled test dataset evaluation dataset that does not influence the training process, as there is no early stopping or feedback loop from the evaluation set. Both training and evaluation use a batch size of 4 per device. The fine-tuning was conducted on a cluster of eight A100 GPUs. It achieves a training speed of approximately 5.84 samples per second on the GPU cluster.

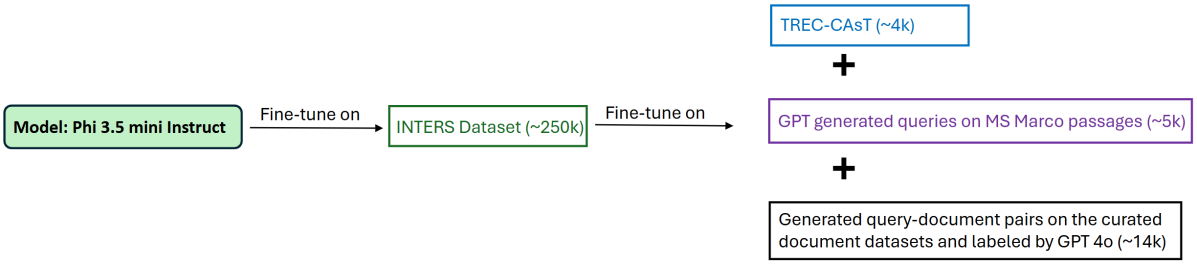

Figure 3: SLM labeler supervised fine-tuning (SFT) workflow: we first fine-tuned Phi-3.5 Mini Instruct on the INTERS dataset for multi-task pre-training, and then fine-tuned on the resulting model again with a combination of synthetic and open-sourced datasets.

Box 3.3: Relevance Labeling Prompt

Role
You are an enterprise search quality rater evaluating file/message relevance.

Task
Given a **query** and an **entity** (document), assign a relevance score in the range of 0–4:
– 4 = ideal quality, should be the ideal result
...
– 0 = bad quality, should never be shown

Input
– Query text
– Document metadata and highlights

Output
`Score: <between 0 and 4>`

4 Experimental Results

In this section, we present a detailed evaluation of the performance of the fine-tuned Phi-3.5 Mini Instruct model compared to the base model (vanilla Phi-3.5 Mini Instruct) and GPT-4o. The results show that fine-tuning significantly enhances the alignment of the Phi-3.5 Mini model with human relevance judgments, matching or even marginally surpassing GPT-4o on both evaluation metrics.

Fine-tuning Impact: Fine-tuning has shown a substantial impact on model performance. As highlighted in Figure 4, the SLM-Human NDCG increased from 0.815 to 0.953, reflecting a significant improvement in the model's alignment with human judgments. Similarly, pairwise accuracy increased from 42.15 to 63.81, surpassing even GPT-4o on both key metrics (NDCG: 0.944, Accuracy: 62.58). This confirms that adapting Phi-3.5 Mini to enterprise relevance tasks yields not only competitive but superior performance, underscoring the effectiveness of our synthetic data generation and fine-tuning pipeline.

We also conducted comprehensive ablation studies on the training data, with results summarized in Table 4. Our key findings are as follows:

Data Size Impact: It is widely known that more training data may yield better results; however, our experiments show that this effect plateaus beyond a certain point. The fine-tuning results highlight the role of data size in shaping model performance. As shown in Table 4, increasing the amount of synthetic training data from 14K to 24K examples (row 2 vs. row 3) yields only marginal gains. With 14K examples, fine-tuning produced a substantial improvement in SLM-Human NDCG (from 0.815 to 0.953) and pairwise accuracy (from 42.15 to 63.81). Expanding to 24K examples, however, resulted in nearly identical performance (NDCG: 0.954, Accuracy: 63.55). A similar pattern is observed when comparing row 4 and row 5, both without multi-task tuning. This suggests that while scaling data initially drives significant improvements, beyond a certain threshold additional synthetic data offers diminishing returns.

Multi-task Tuning: Multi-task tuning demonstrated a notable improvement in aligning model outputs with human relevance judgments. As shown in Table 4, the impact of multi-task tuning becomes clear when comparing row 2 (with multi-task) against row 4 (without). Under the same 14K synthetic training dataset, the model with multi-task tuning achieved NDCG of 0.953 and accuracy of 63.81, compared to 0.946 and 62.41 without multi-task. This increase suggests that multi-task tuning, in conjunction with instruct tuning, effectively enhances the model's ability to generalize to human relevance preferences. A similar pattern holds when comparing row 3 (with multi-task) and row 5 (without) under the 24K synthetic dataset, where NDCG improves from 0.945 to 0.954 and accuracy from 62.70 to 63.55.

To further validate the role of our synthetic queries, we fine-tuned the model without including the synthetic enterprise query dataset, while keeping all other open-source datasets unchanged (i.e., excluding the data in the black box in Figure 3). As shown in row 7 of Table 4, this variant yields only marginal change over the vanilla model, with NDCG moving from 0.815 to 0.826 and Accuracy from 42.16 to 42.88. This outcome aligns with our expectation: existing open-source datasets are primarily designed for web search query understanding and thus do not solely capture the characteristics of enterprise search relevance labeling. These results confirm that our synthetic enterprise queries provide the primary training signal, while multi-task tuning offers an auxiliary benefit by improving robustness via multi-task tuning.

Training Dataset Quality: As detailed in Subsection 3.1, we refine the generated queries from the LLM through a query refinement step to enhance their credibility and diversity. When this refinement was applied after synthetic query generation, the model (row 2 in Table 4) achieved strong performance with NDCG of 0.953 and accuracy of 63.81. In contrast, reusing the same 14K raw synthetic queries without refinement (row 6) led to a measurable drop (NDCG 0.943, Accuracy 60.97). This underscores the critical importance of training data quality. Moreover, increasing the dataset size from 14K to 24K examples (row 2 vs. row 3) produced almost identical results, highlighting that for fine-tuning the SLM, improving data quality has a greater impact than simply adding more data, once the dataset surpasses a reasonably large threshold.

Throughput Analysis: Our fine-tuned SLM achieved 873.33 RPM on a single A100 GPU, which extrapolates to nearly 7K RPM on a cluster of eight A100s. This throughput is an order of magnitude faster than typical LLM labelers, which are generally limited to the hundreds of RPM range.

Cost Analysis: Cost is a major consideration, since large-scale evaluation and training data labeling are both frequent and expensive. Our fine-tuned Phi-3.5 Mini Instruct offers a substantial advantage over large language models such as GPT-4o, being nearly 19× cheaper on both input and output tokens. Specifically, the cost is \$0.13 vs. \$2.50 per 1M input tokens and \$0.52 vs. \$10.00 per 1M output tokens [48]. This cost efficiency, combined with strong alignment to human judgments, makes trained SLM an attractive option for scalable enterprise relevance labeling.

Statistical Significance Analysis To further validate the effectiveness of our fine-tuned labeler, we conducted statistical significance testing against the GPT-4o baseline. Our benchmark dataset consists of 923 query–document pairs with human annotations, spanning 228 distinct queries used for evaluation. For each query, we computed the pairwise accuracy and NDCG of both our fine-tuned labeler and GPT-4o with respect to human annotations. Let $\delta_q = \text{labeler}_q - \text{GPT-4o}_q$ denote the per-query performance difference.

Since our goal is to demonstrate that our labeler is on par with the GPT-4o labeler with significantly less throughput and cost, we performed a one-sided paired non-inferiority test using the Wilcoxon signed-rank test [49]. The null hypothesis assumes that our model performs worse than the baseline by more than a predefined non-inferiority margin Δ, i.e.,

$$H_0 : \mathbb{E}(\delta_q) \leq -\Delta, \qquad H_1 : \mathbb{E}(\delta_q) > -\Delta.$$

We adopt conservative non-inferiority margins of $\Delta = 0.1\%$ for pairwise accuracy and $\Delta = 0.0001$ for NDCG, which represent negligible differences relative to the scale of the respective metrics.

The resulting one-sided Wilcoxon signed-rank tests yield p-values of 0.012 for pairwise accuracy and 0.00098 for NDCG, both of which are below the 0.05 significance level. Therefore, we reject the null hypothesis for both metrics, indicating that our fine-tuned labeler is non-inferior to the GPT-4o baseline within the specified small margins. These results provide statistical evidence, from both pairwise accuracy and NDCG, that our fine-tuned model achieves performance comparable to the strong GPT-4o labeler.

Qualitative Analysis: After examining examples from the human-annotated evaluation set, we observe a clear improvement in our fine-tuned labeler over the vanilla SLM. In particular, the original SLM and even GPT-4o tend to avoid assigning the lowest relevance score, even for documents that are completely off-topic, leading to overly optimistic relevance judgments. By contrast, our fine-tuned query labeler demonstrates greater calibration and confidence: it is

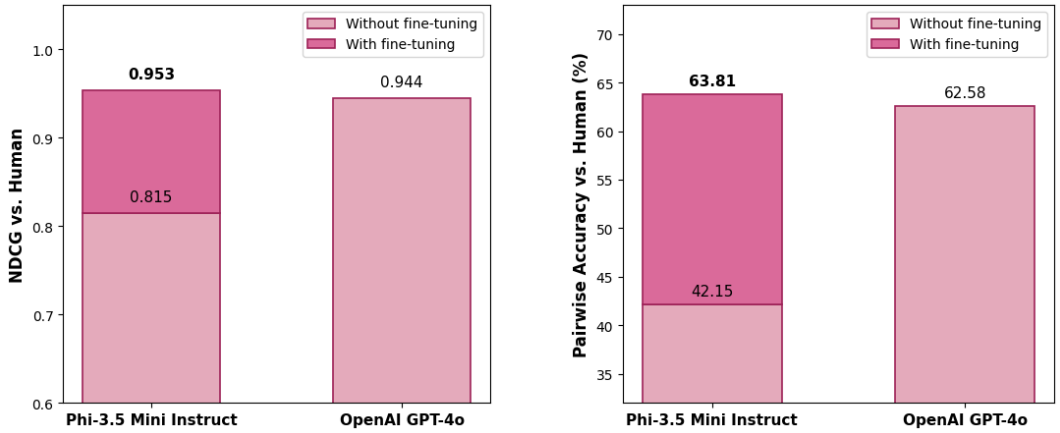

Figure 4: Performance results comparing the vanilla and fine-tuned SLM models against GPT-4o on two key evaluation metrics for the enterprise search test dataset.

Row	Multi-task Tuning	Data Revision	Synthetic Data Size	SLM-Human		SLM-LLM		LLM-Human	
				NDCG	Accuracy (%)	NDCG	Accuracy (%)	NDCG	Accuracy (%)
1	–	–	– (Vanilla SLM)	0.815	42.15	0.820	41.38	0.944	62.58
2	✓	✓	14K	0.953	63.81	0.951	66.16		
3	✓	✓	24K	0.954	63.55	0.950	66.10		
4	✗	✓	14K	0.946	62.41	0.950	66.72		
5	✗	✓	24K	0.943	62.50	0.948	66.83		
6	✗	✗	14K	0.943	60.97	0.948	66.05		
7	✓	✓	0 (Only Public Data)	0.826	42.88	0.824	42.19		

Table 4: Performance results comparing the vanilla and trained SLMs under different fine-tuning configurations for ablation studies. The last three column groups present the pairwise metrics based on labels from SLM/SLM/LLM against Human/LLM/Human.

able to correctly assign a score of 0 when a passage is entirely irrelevant, thereby aligning more closely with human annotations.

These results collectively demonstrate that carefully fine-tuning SLMs with high-quality synthetic data not only achieves a level of alignment with human judgments comparable to that of large LLMs, but also enables substantially more scalable and cost-efficient relevance labeling.

5 Conclusion

Our work presents an efficient and scalable approach to enterprise search relevance labeling by fine-tuning Small Language Models (SLMs) with a synthetic data generation pipeline. The fine-tuned Phi-3.5 Mini Instruct achieves GPT-4o–level accuracy while significantly improving throughput and reducing costs, making large-scale query–document relevance labeling practical. By leveraging GPT-4o for query generation, BM25 for hard negative mining, and LLM-based judgments for labeling, we create high-quality synthetic training data that enables SLMs to perform competitively with LLMs.

Empirical results show our fine-tuned model achieves SLM Human NDCG of 0.953 and pairwise accuracy of 63.81, outperforming GPT-4o (NDCG: 0.944, Accuracy: 62.58). These findings highlight that compact, well-trained SLMs can match or even outperform large models in aligning with human relevance judgments, while enabling faster and more sustainable annotation pipelines. More broadly, this work highlights the potential of synthetic data–driven training to bridge the gap between high-quality human-aligned evaluation and scalable model development, offering a practical path toward deploying lightweight models for data-intensive NLP tasks.

References

[1] Yutao Zhu, Huaying Yuan, Shuting Wang, Jiongnan Liu, Wenhan Liu, Chenlong Deng, Haonan Chen, Zheng Liu, Zhicheng Dou, and Ji-Rong Wen. Large language models for information retrieval: A survey, 2024.

[2] Cunxiang Wang, Xiaoze Liu, Yuanhao Yue, Xiangru Tang, Tianhang Zhang, Cheng Jiayang, Yunzhi Yao, Wenyang Gao, Xuming Hu, Zehan Qi, Yidong Wang, Linyi Yang, Jindong Wang, Xing Xie, Zheng Zhang, and Yue Zhang. Survey on factuality in large language models: Knowledge, retrieval and domain-specificity, 2023.

[3] David Hawking. Challenges in enterprise search. In *ADC*, volume 4, pages 15–24, 2004.

[4] Filip Radlinski and Nick Craswell. A theoretical framework for conversational search. In *Proceedings of the 2017 conference on conference human information interaction and retrieval*, pages 117–126, 2017.

[5] Tri Nguyen, Mir Rosenberg, Xia Song, Jianfeng Gao, Saurabh Tiwary, Rangan Majumder, and Li Deng. Ms marco: A human-generated machine reading comprehension dataset. 2016.

[6] Jeffrey Dalton, Chenyan Xiong, and Jamie Callan. Trec cast 2019: The conversational assistance track overview. *arXiv preprint arXiv:2003.13624*, 2020.

[7] Nayoung Choi, Youngjune Lee, Gyu-Hwung Cho, Haeyu Jeong, Jungmin Kong, Saehun Kim, Keunchan Park, Sarah Cho, Inchang Jeong, Gyohee Nam, et al. Rradistill: Distilling llms' passage ranking ability for long-tail queries document re-ranking on a search engine. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing: Industry Track*, pages 627–641, 2024.

[8] Paul Thomas, Seth Spielman, Nick Craswell, and Bhaskar Mitra. Large language models can accurately predict searcher preferences. In *Proceedings of the 47th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 1930–1940, 2024.

[9] Marah Abdin, Jyoti Aneja, Hany Awadalla, Ahmed Awadallah, Ammar Ahmad Awan, and et al. Phi-3 technical report: A highly capable language model locally on your phone, 2024.

[10] Quentin Fitte-Rey, Matyas Amrouche, and Romain Deveaud. Augmented relevance datasets with fine-tuned small llms. *arXiv preprint arXiv:2504.09816*, 2025.

[11] Stephen E Robertson, Steve Walker, Susan Jones, Micheline M Hancock-Beaulieu, Mike Gatford, et al. *Okapi at TREC-3*. British Library Research and Development Department, 1995.

[12] Marah Abdin, Jyoti Aneja, Hany Awadalla, Ahmed Awadallah, Ammar Ahmad Awan, Nguyen Bach, Amit Bahree, Arash Bakhtiari, Jianmin Bao, Harkirat Behl, et al. Phi-3 technical report: A highly capable language model locally on your phone. *arXiv preprint arXiv:2404.14219*, 2024.

[13] Humza Naveed, Asad Ullah Khan, Shi Qiu, Muhammad Saqib, Saeed Anwar, Muhammad Usman, Naveed Akhtar, Nick Barnes, and Ajmal Mian. A comprehensive overview of large language models. *ACM Transactions on Intelligent Systems and Technology*, 16(5):1–72, 2025.

[14] Shijie Wu, Ozan Irsoy, Steven Lu, Vadim Dabravolski, Mark Dredze, Sebastian Gehrmann, Prabhanjan Kambadur, David Rosenberg, and Gideon Mann. Bloomberggpt: A large language model for finance, 2023. *URL https://arxiv.org/abs/2303.17564*, 2024.

[15] Arun James Thirunavukarasu, Darren Shu Jeng Ting, Kabilan Elangovan, Laura Gutierrez, Ting Fang Tan, and Daniel Shu Wei Ting. Large language models in medicine. *Nature medicine*, 29(8):1930–1940, 2023.

[16] Matthew E. Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. Deep contextualized word representations. In Marilyn Walker, Heng Ji, and Amanda Stent, editors, *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 2227–2237, New Orleans, Louisiana, June 2018. Association for Computational Linguistics.

[17] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 conference of the North American chapter of the association for computational linguistics: human language technologies, volume 1 (long and short papers)*, pages 4171–4186, 2019.

[18] Kevin Clark, Minh-Thang Luong, Quoc V. Le, and Christopher D. Manning. ELECTRA: Pre-training text encoders as discriminators rather than generators. In *ICLR*, 2020.

[19] Pengcheng He, Xiaodong Liu, Jianfeng Gao, and Weizhu Chen. Deberta: Decoding-enhanced bert with disentangled attention. In *International Conference on Learning Representations*, 2021.

[20] Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *Journal of machine learning research*, 21(140):1–67, 2020.

[21] Omar Khattab and Matei Zaharia. Colbert: Efficient and effective passage search via contextualized late interaction over bert. In *Proceedings of the 43rd International ACM SIGIR conference on research and development in Information Retrieval*, pages 39–48, 2020.

[22] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.

[23] Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Ostrow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card. *arXiv preprint arXiv:2410.21276*, 2024.

[24] Yutao Zhu, Huaying Yuan, Shuting Wang, Jiongnan Liu, Wenhan Liu, Chenlong Deng, Haonan Chen, Zheng Liu, Zhicheng Dou, and Ji-Rong Wen. Large language models for information retrieval: A survey. *arXiv preprint arXiv:2308.07107*, 2023.

[25] Gemma Team, Morgane Riviere, Shreya Pathak, Pier Giuseppe Sessa, Cassidy Hardin, Surya Bhupatiraju, Léonard Hussenot, Thomas Mesnard, Bobak Shahriari, Alexandre Ramé, et al. Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118*, 2024.

[26] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.

[27] Chris Samarinas and Hamed Zamani. Distillation and refinement of reasoning in small language models for document re-ranking. In *Proceedings of the 2025 International ACM SIGIR Conference on Innovative Concepts and Theories in Information Retrieval (ICTIR)*, pages 430–435, 2025.

[28] Sean MacAvaney and Luca Soldaini. One-shot labeling for automatic relevance estimation. In *Proceedings of the 46th International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 2230–2235, 2023.

[29] Zahra Abbasiantaeb, Chuan Meng, Leif Azzopardi, and Mohammad Aliannejadi. Can we use large language models to fill relevance judgment holes? *arXiv preprint arXiv:2405.05600*, 2024.

[30] Naghmeh Farzi and Laura Dietz. Criteria-based llm relevance judgments. In *Proceedings of the 2025 International ACM SIGIR Conference on Innovative Concepts and Theories in Information Retrieval (ICTIR)*, pages 254–263, 2025.

[31] Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, Weizhu Chen, et al. Lora: Low-rank adaptation of large language models. *ICLR*, 1(2):3, 2022.

[32] Navid Mehrdad, Hrushikesh Mohapatra, Mossaab Bagdouri, Prijith Chandran, Alessandro Magnani, Xunfan Cai, Ajit Puthenputhussery, Sachin Yadav, Tony Lee, ChengXiang Zhai, et al. Large language models for relevance judgment in product search. *arXiv preprint arXiv:2406.00247*, 2024.

[33] Orion Weller, Kathryn Ricci, Eugene Yang, Andrew Yates, Dawn Lawrie, and Benjamin Van Durme. Rank1: Test-time compute for reranking in information retrieval. *arXiv preprint arXiv:2502.18418*, 2025.

[34] Aaron Jaech, Adam Kalai, Adam Lerer, Adam Richardson, Ahmed El-Kishky, Aiden Low, Alec Helyar, Aleksander Madry, Alex Beutel, Alex Carney, et al. Openai o1 system card. *arXiv preprint arXiv:2412.16720*, 2024.

[35] Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.

[36] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837, 2022.

[37] Marco Braga, Pranav Kasela, Alessandro Raganato, and Gabriella Pasi. Synthetic data generation with large language models for personalized community question answering. In *2024 IEEE/WIC International Conference on Web Intelligence and Intelligent Agent Technology (WI-IAT)*, pages 360–366, 2024.

[38] Luiz Bonifacio, Hugo Abonizio, Marzieh Fadaee, and Rodrigo Nogueira. InPars: Unsupervised dataset generation for information retrieval. In *Proceedings of the 45th International ACM SIGIR Conference on Research and Development in Information Retrieval*, SIGIR '22, page 2387–2392, New York, NY, USA, 2022. Association for Computing Machinery.

[39] Zhuyun Dai, Vincent Y Zhao, Ji Ma, Yi Luan, Jianmo Ni, Jing Lu, Anton Bakalov, Kelvin Guu, Keith Hall, and Ming-Wei Chang. Promptagator: Few-shot dense retrieval from 8 examples. In *The Eleventh International Conference on Learning Representations*, 2023.

[40] Vitor Jeronymo, Luiz Bonifacio, Hugo Abonizio, Marzieh Fadaee, Roberto Lotufo, Jakub Zavrel, and Rodrigo Nogueira. InPars-v2: Large language models as efficient dataset generators for information retrieval, 2023.

[41] Ramraj Chandradevan, Kaustubh Dhole, and Eugene Agichtein. DUQGen: Effective unsupervised domain adaptation of neural rankers by diversifying synthetic query generation. In Kevin Duh, Helena Gomez, and Steven Bethard, editors, *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 7430–7444, Mexico City, Mexico, June 2024. Association for Computational Linguistics.

[42] Yutao Zhu, Peitian Zhang, Chenghao Zhang, Yifei Chen, Binyu Xie, Zheng Liu, Ji-Rong Wen, and Zhicheng Dou. INTERS: Unlocking the power of large language models in search with instruction tuning. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2782–2809, Bangkok, Thailand, August 2024. Association for Computational Linguistics.

[43] Yining Wang, Liwei Wang, Yuanzhi Li, Di He, and Tie-Yan Liu. A theoretical analysis of ndcg type ranking measures. In *Conference on learning theory*, pages 25–54. PMLR, 2013.

[44] Ali Vardasbi, Harrie Oosterhuis, and Maarten de Rijke. When inverse propensity scoring does not work: Affine corrections for unbiased learning to rank. In *Proceedings of the 29th ACM International Conference on Information & Knowledge Management*, pages 1475–1484, 2020.

[45] Nick Craswell, Onno Zoeter, Michael Taylor, and Bill Ramsey. An experimental comparison of click position-bias models. In *Proceedings of the 2008 international conference on web search and data mining*, pages 87–94, 2008.

[46] Stephen Robertson, Hugo Zaragoza, et al. The probabilistic relevance framework: Bm25 and beyond. *Foundations and Trends® in Information Retrieval*, 3(4):333–389, 2009.

[47] Shengyu Zhang, Linfeng Dong, Xiaoya Li, Sen Zhang, Xiaofei Sun, Shuhe Wang, Jiwei Li, Runyi Hu, Tianwei Zhang, Fei Wu, et al. Instruction tuning for large language models: A survey. *arXiv preprint arXiv:2308.10792*, 2023.

[48] OpenAI API Pricing. `https://openai.com/api/pricing/`, 2025.

[49] Frank Wilcoxon. Individual comparisons by ranking methods. In *Breakthroughs in statistics: Methodology and distribution*, pages 196–202. Springer, 1992.